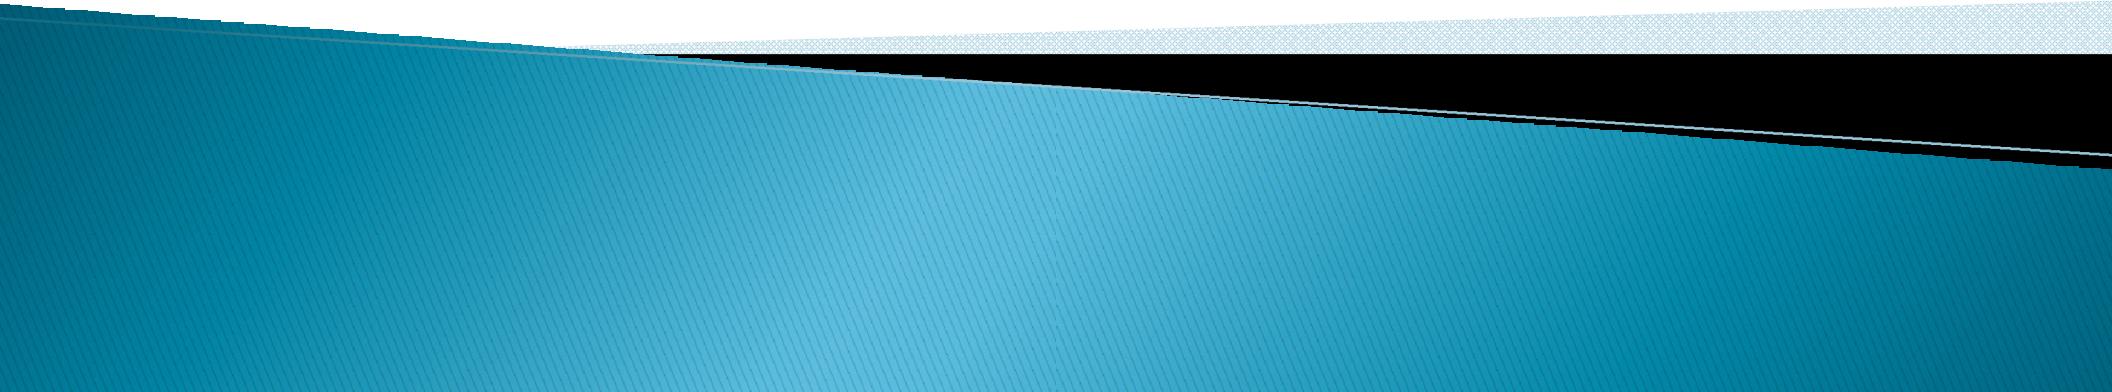


Apache Kafka





Why Kafka?

When we have....
Aren't they Good?



*Apache ActiveMQ, JBoss HornetQ, Zero MQ, RabbitMQ are respective brands of Apache Software Foundation, JBoss Inc, iMatix Corporation, Vmware Inc.

They all are Good
But not for all use-cases.



Use Cases

Real-Time Stream Processing (combined with Spark Streaming)

General purpose Message Bus

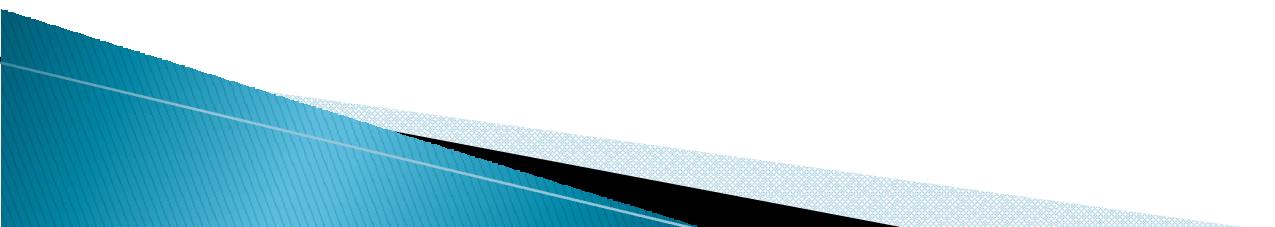
Collecting User Activity Data

Collecting Operational Metrics from applications, servers or devices

Log Aggregation

Change Data Capture

Commit Log for distributed systems



What is Common?

Scalable: Need to be Highly Scalable. A lot of Data. It can be billions of message.

Reliability of messages, What If, I loose a small no. of messages. Is it **fine** with me ?.

Distributed : Multiple Producers, Multiple Consumers

High-throughput: Does not require to have **JMS** Standards, as it may be overkill for some use-cases like transportation of logs.

- As per JMS, each message has to be **acknowledged** back.
- Exactly one delivery guarantee requires **two-phase commit**.



Introduction

An Apache Project, initially developed by LinkedIn's SNA team.
A High-throughput distributed Publish-Subscribe based messaging system.

A Kind of Data Pipeline

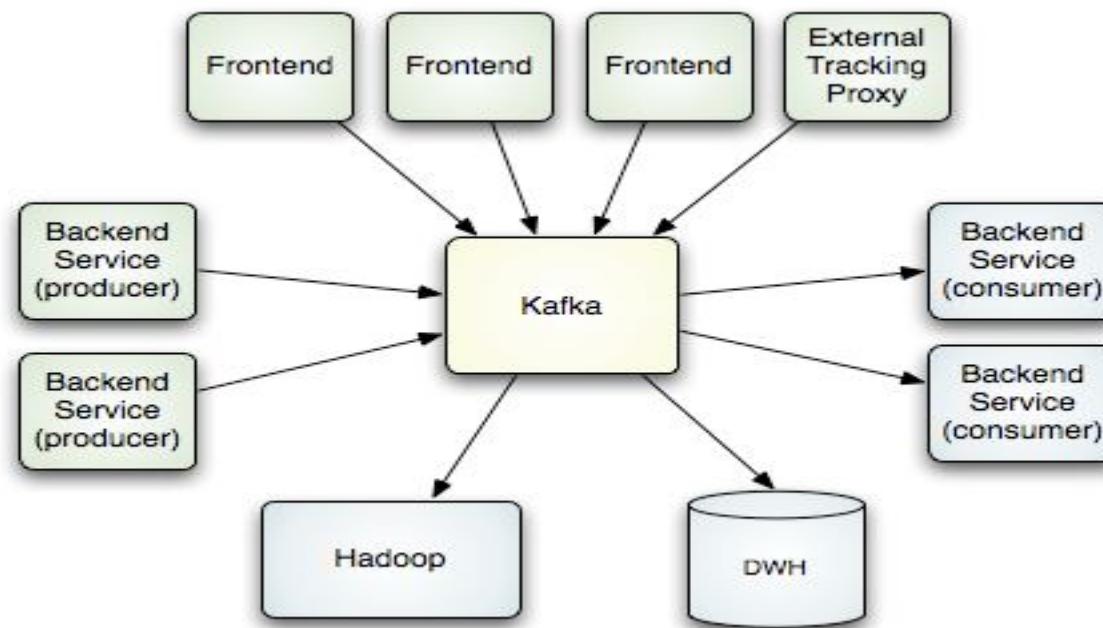
Written in Scala.

Does not follow JMS Standards, neither uses JMS APIs.

Supports both queue and topic semantics.



Overview



Credit : <http://kafka.apache.org/design.html>

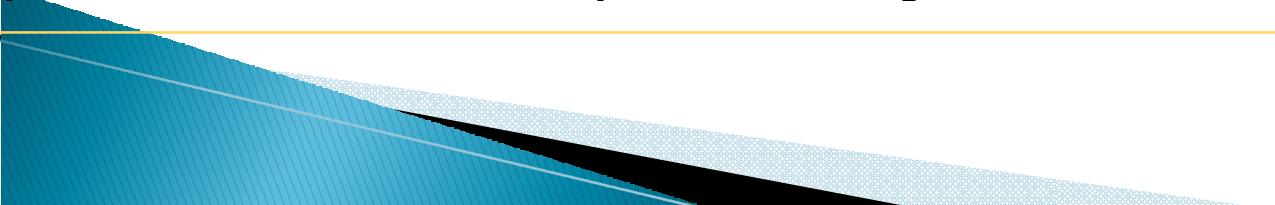
Key terminology

Kafka maintains feeds of messages in categories called *topics*.
Processes that publish messages to a Kafka topic are called *producers*.

Processes that subscribe to topics and process the feed of published messages are called *consumers*.

Kafka is run as a cluster comprised of one or more servers each of which is called a *broker*.

Communication between all components is done via a high performance simple binary API over TCP protocol

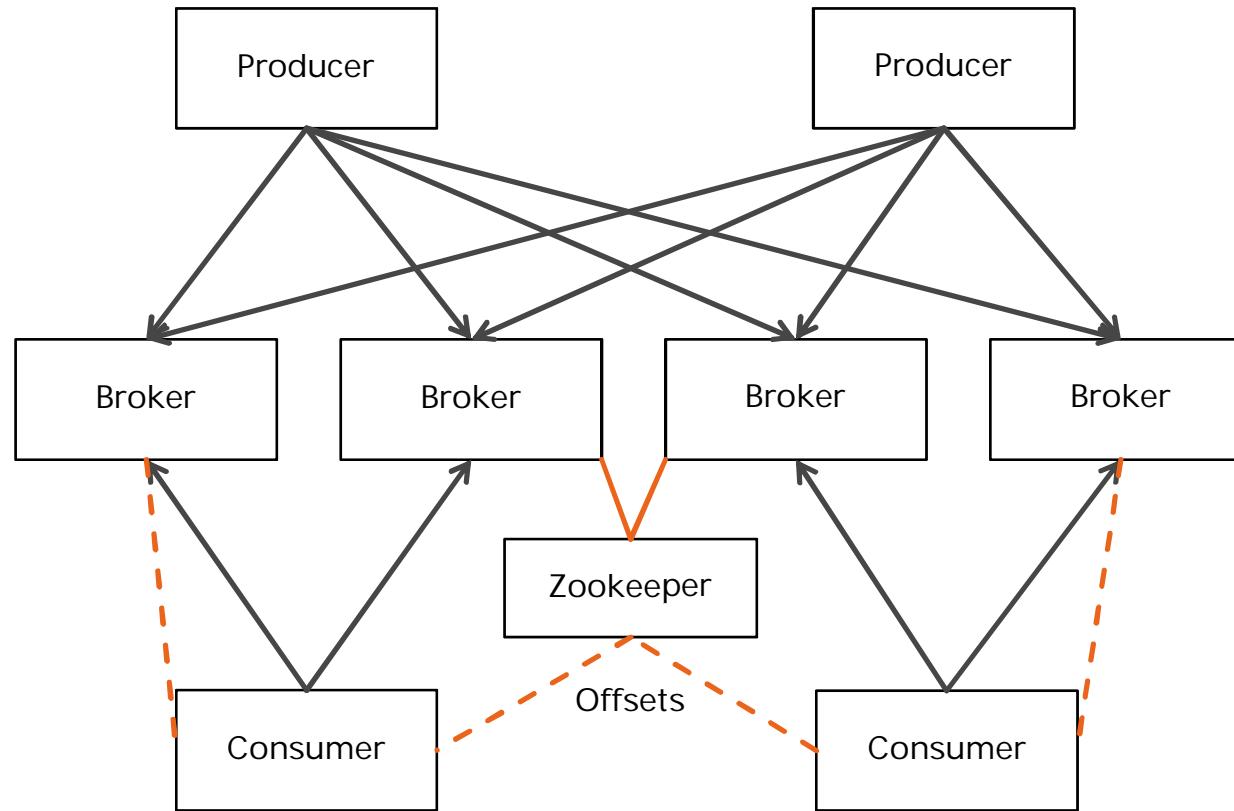


Architecture

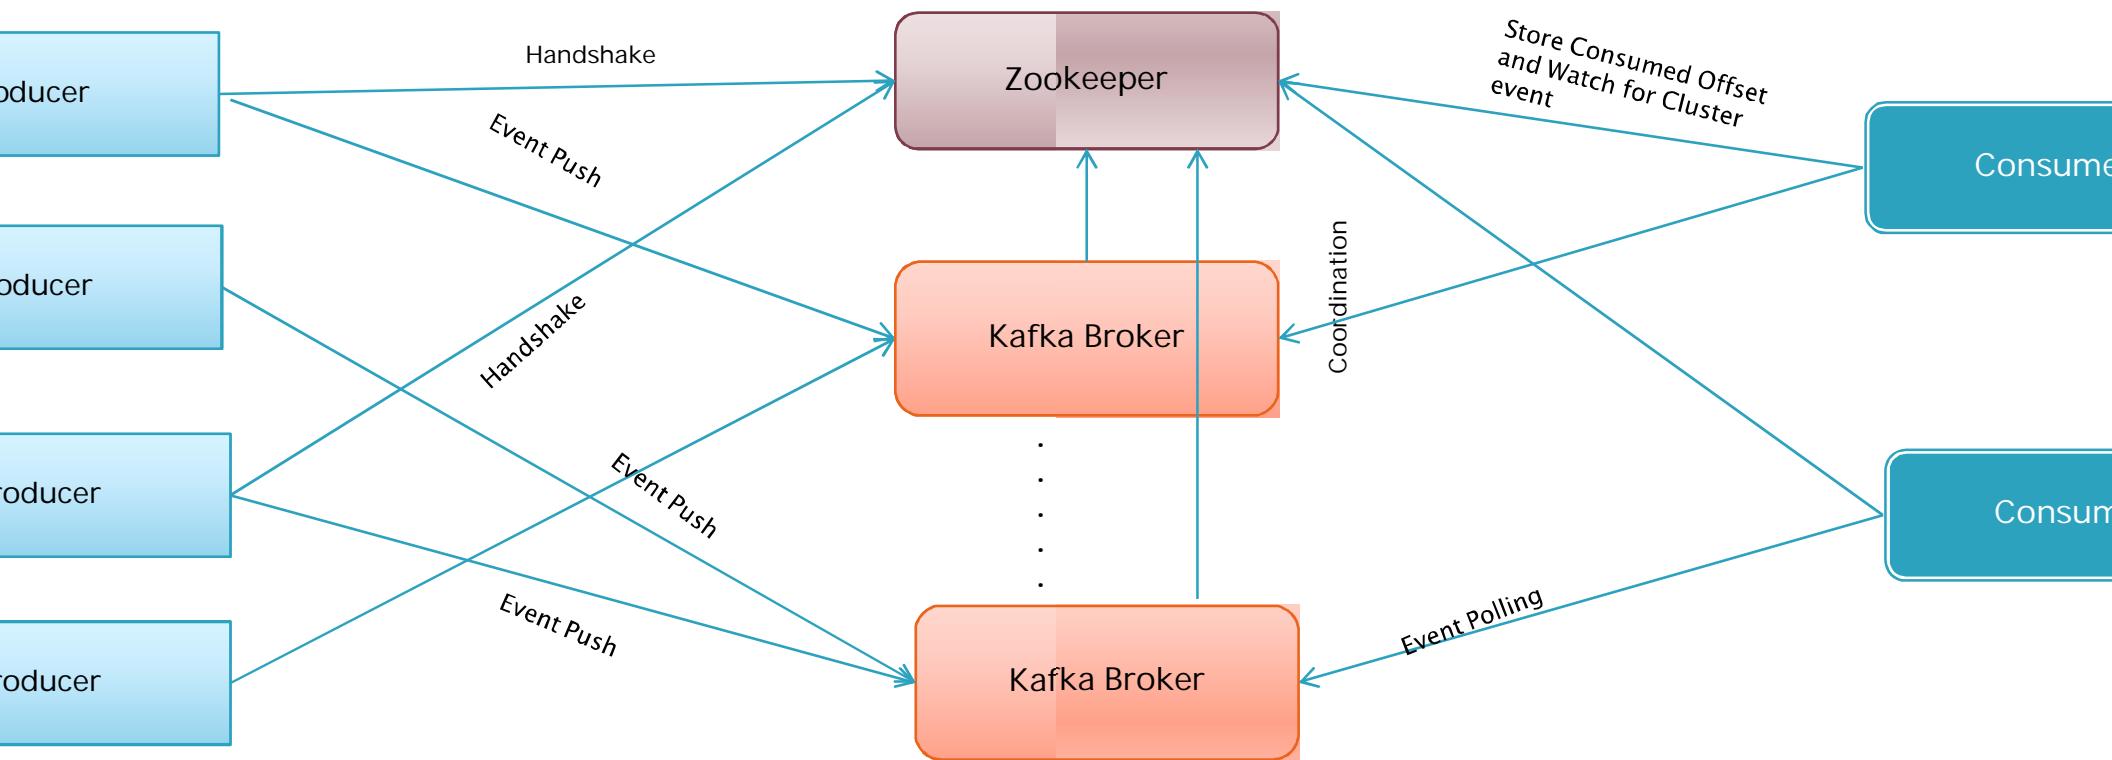
Producers

Zookeeper

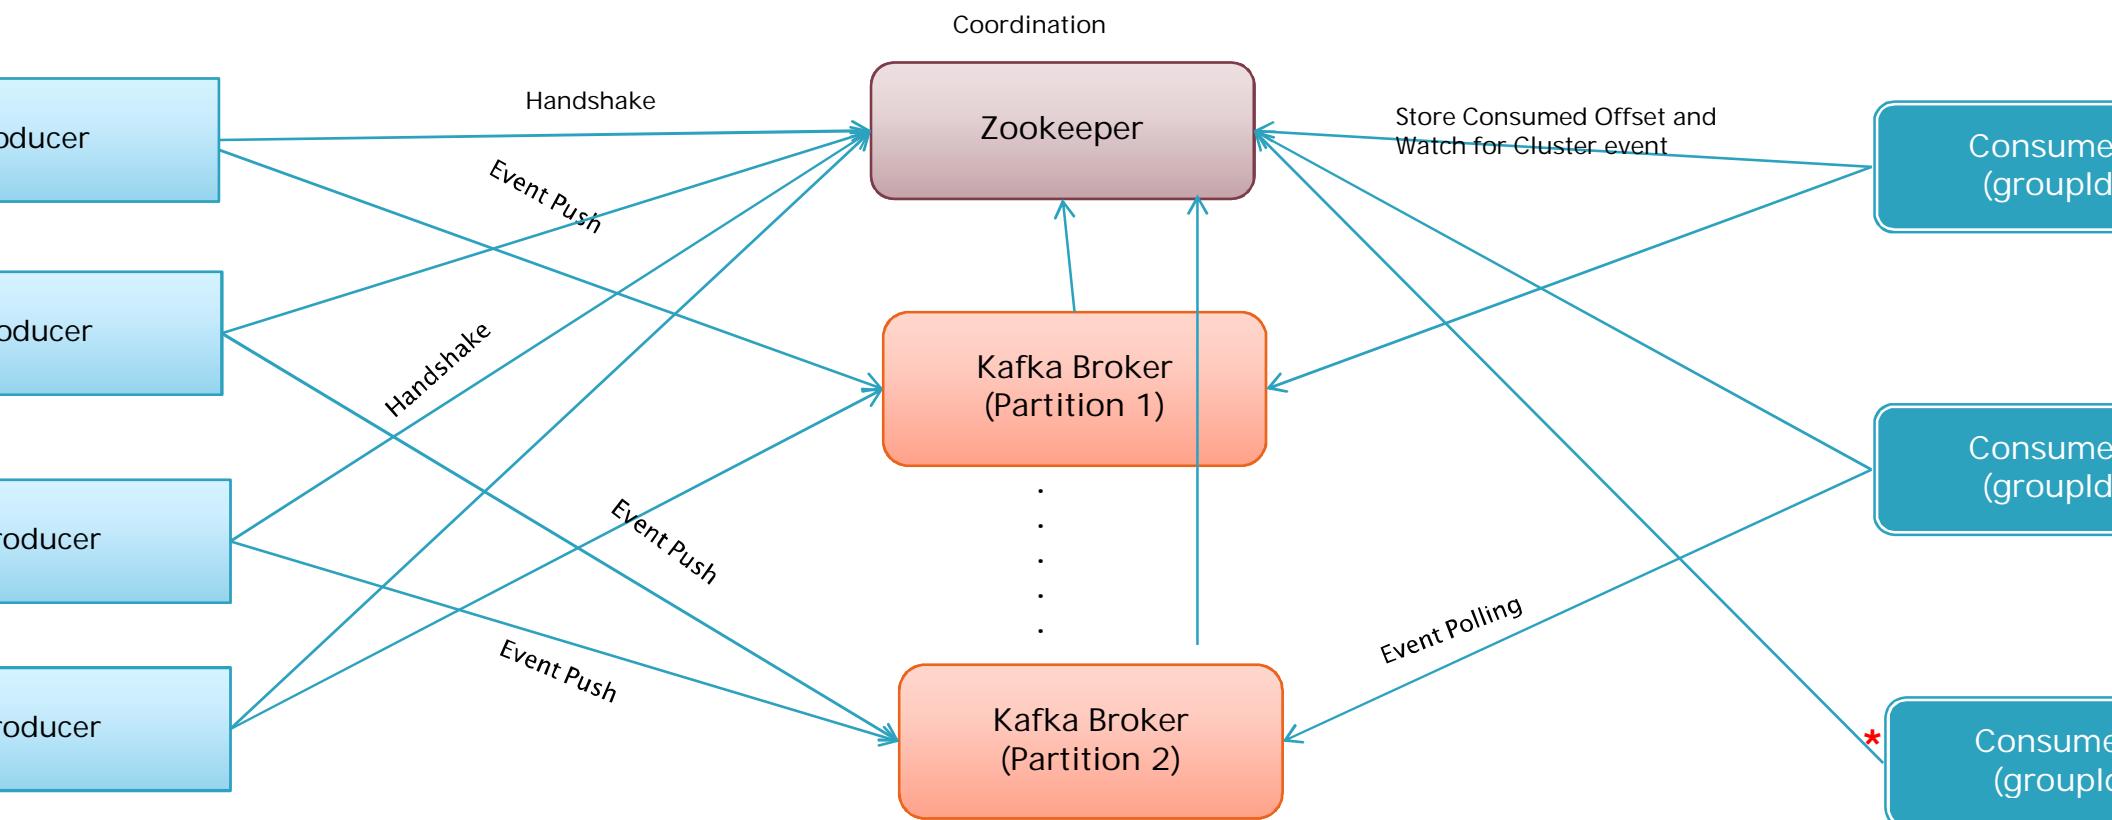
Consumers



How it works



How it works (Queue)



* Consumer 3 would not receive any data, as number of consumers are more than number of partitions.

Design Elements

Filesystem Cache

Zero-copy transfer of messages

Batching of Messages

Batch Compression

Automatic Producer Load balancing.

Broker does not **Push** messages to Consumer, Consumer **Poll** messages from Broker.

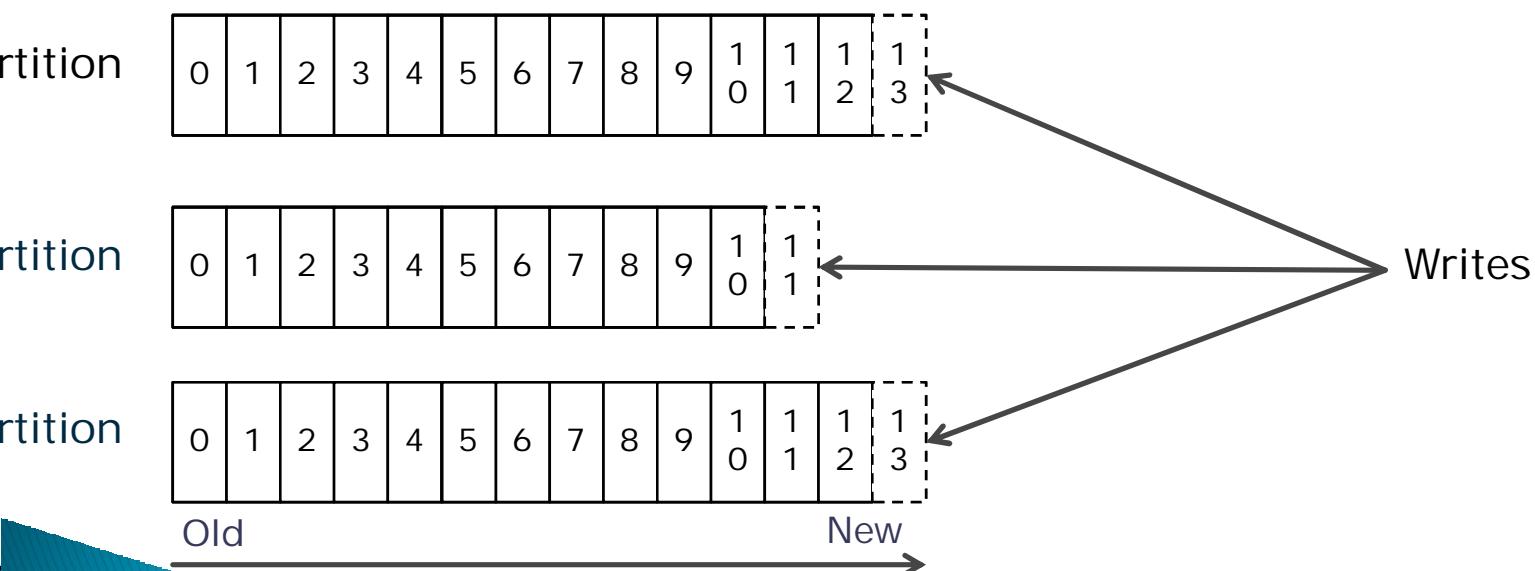
And Some others.

- Cluster formation of Broker/Consumer using Zookeeper, So on the fly more consumer, broker can be introduced. The new cluster rebalancing will be taken care by Zookeeper
- Data is persisted in broker and is not removed on consumption (till retention period), so if one consumer fails while consuming, same message can be re-consume again later from broker.
- Simplified storage mechanism for message, not for each message per consumer.

Topics – Partitions

Topics are broken up into ordered commit logs called partitions.
Each message in a partition is assigned a sequential id called an offset.

Data is retained for a configurable period of time*



Message Ordering

Ordering is only guaranteed within a partition for a topic
To ensure ordering:

- Group messages in a partition by key (producer)
- Configure exactly one consumer instance per partition within a consumer group



Guarantees

Messages sent by a producer to a particular topic partition will be appended in the order they are sent

A consumer instance sees messages in the order they are stored in the log

For a topic with replication factor N, Kafka can tolerate up to $N-1$ server failures without “losing” any messages committed to the log



Topics – Replication

Topics can (and should) be replicated.

The unit of replication is the partition

Each partition in a topic has 1 leader and 0 or more replicas.

A replica is deemed to be “in-sync” if

- The replica can communicate with Zookeeper
- The replica is not “too far” behind the leader (configurable)

The group of in-sync replicas for a partition is called the **ISR**
(In-Sync Replicas)

The Replication factor cannot be lowered



Topics – Replication

Durability can be configured with the producer configuration
request.required.acks

- 0** The producer never waits for an ack
- 1** The producer gets an ack after the leader replica has received the data
- 1** The producer gets an ack after all ISRs receive the data

The minimum available ISR can also be configured such that an error is returned if enough replicas are not available to replicate data.



Durable Writes

producers can choose to *trade throughput for durability of writes*:

Reliability	Behaviour	Per Event Latency	Required Acknowledgements (request.required.acks)
Highest	ACK all ISRs have received	Highest	-1
Medium	ACK once the leader has received	Medium	1
Lowest	No ACKs required	Lowest	0

Throughput can also be raised with ***more brokers***... (so do this instead)!

A sane configuration:

Property	Value
replication	3
min.insync.replicas	2
request.required.acks	-1

Producer

producers publish to a topic of their choosing (push)
load can be distributed

Typically by “round-robin”

Can also do “semantic partitioning” based on a key in the message
brokers load balance by partition

can support async (less durable) sending

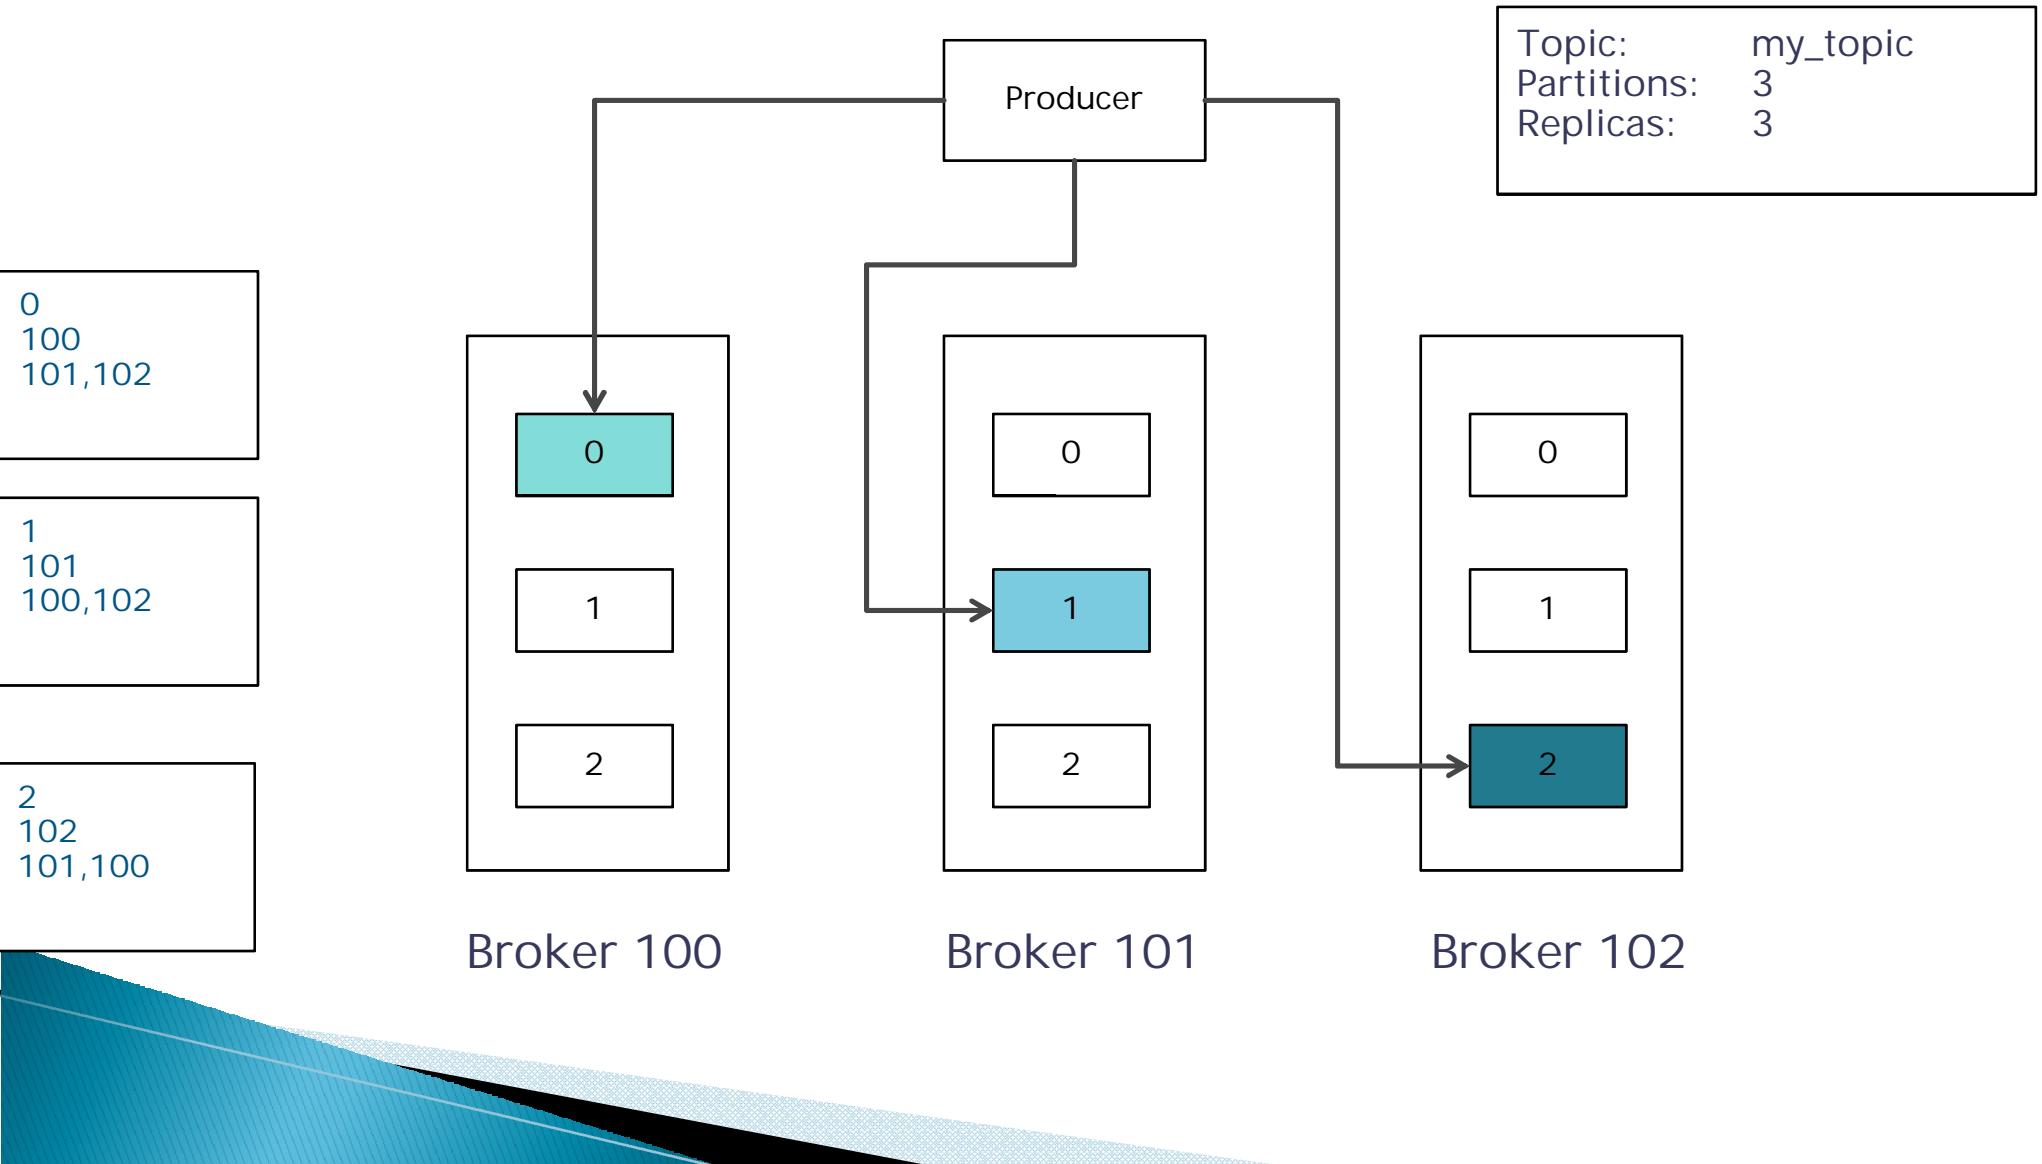
All nodes can answer metadata requests about:

Which servers are alive

Where leaders are for the partitions of a topic



Producer - Load Balancing and ISRs



Consumer – Groups

Consumers can be organized into Consumer Groups
common Patterns:

1) All consumer instances in one group

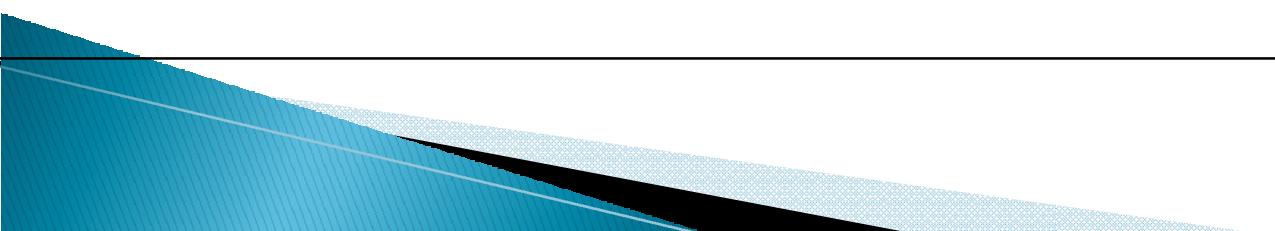
- Acts like a traditional queue with load balancing

2) All consumer instances in different groups

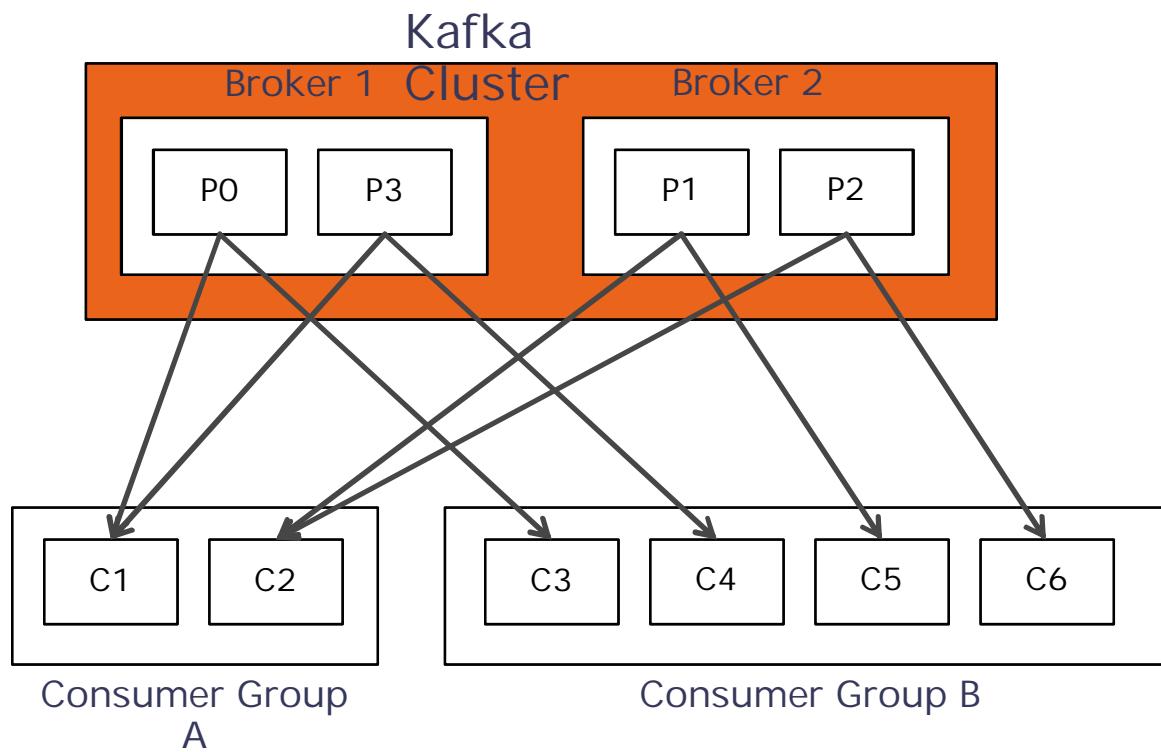
- All messages are broadcast to all consumer instances

3) “Logical Subscriber” – Many consumer instances in a group

- Consumers are added for scalability and fault tolerance
- Each consumer instance reads from one or more partitions for a topic
- There cannot be more consumer instances than partitions

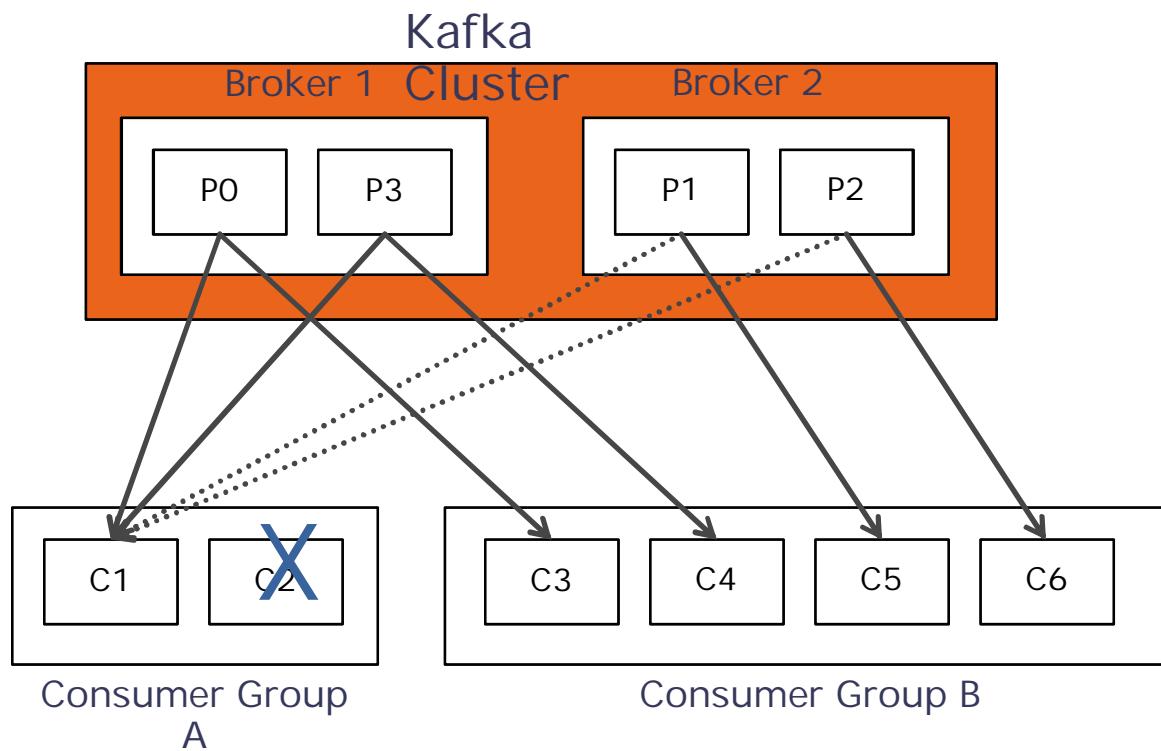


Consumer – Groups



Consumer
Groups provide
convenience and
scalability

Consumer – Groups



rebalance
themselves

Delivery Semantics

Default

At least once

- **Messages are never lost but may be redelivered**

At most once

- Messages are lost but never redelivered

Exactly once

- Messages are delivered once and only once

Delivery Semantics

At least once

- **Messages are never lost but may be redelivered**

At most once

- Messages are lost but never redelivered

Exactly once

- Messages are delivered once and only once

Much Harder
(impossible??)

Getting Exactly Once Semantics

Must consider two components

- Durability guarantees when *publishing* a message
- Durability guarantees when *consuming* a message

Producer

- What happens when a produce request was sent but a network error returned before an ack?
- Use a single writer per partition and check the latest committed value after network errors

Consumer

- Include a unique ID (e.g. UUID) and de-duplicate.
- Consider storing offsets with data



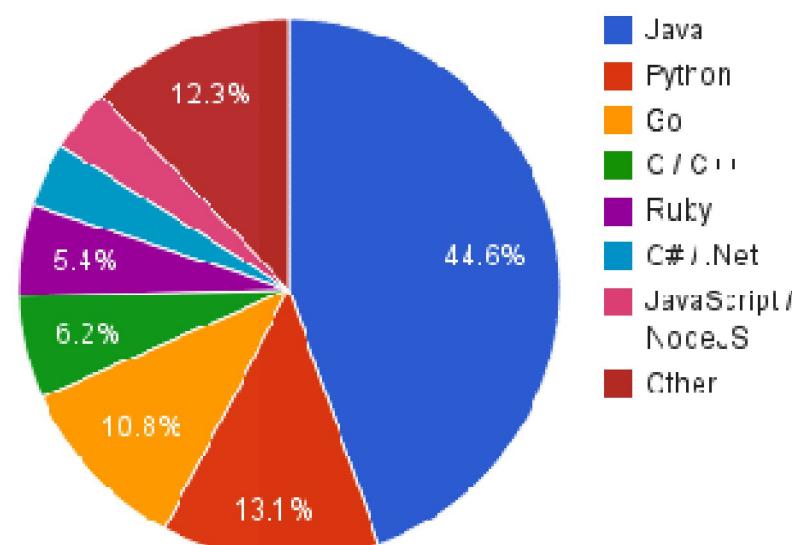
Kafka Clients

Remember Kafka implements a binary TCP protocol.

All clients except the JVM-based clients are maintained external to the code base.

Full Client List [here](#)

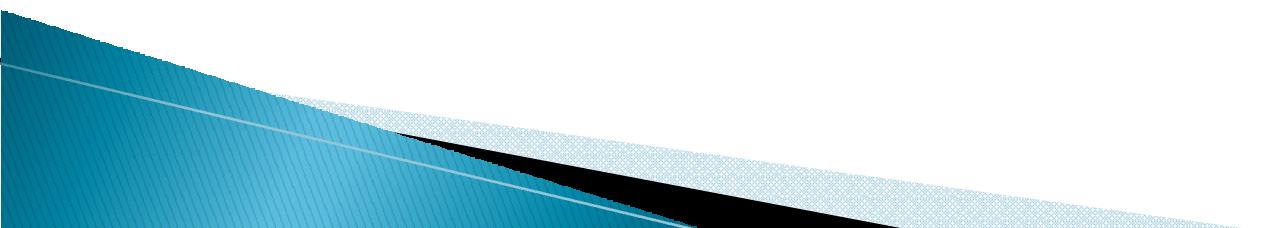
Kafka Producer/Consumer Language



Java Producer Example - Old (< 0.8.1)

```
/* start the producer */
private void start() {
producer = new Producer<String, String>](config);
}

/* create record and send to Kafka
 * because the key is null, data will be sent to a random partition.
 * the producer will switch to a different random partition every 10
tes
*/
private void produce(String s) {
    KeyedMessage<String, String> message = new
dMessage<String, String>](topic, null, s);
    producer.send(message);
}
```



Producer – New

Send the given record asynchronously and return a future which will eventually contain the response information.

@param record The record to send

@return A future which will eventually contain the response information

public Future<RecordMetadata> send(ProducerRecord<K, V> record);

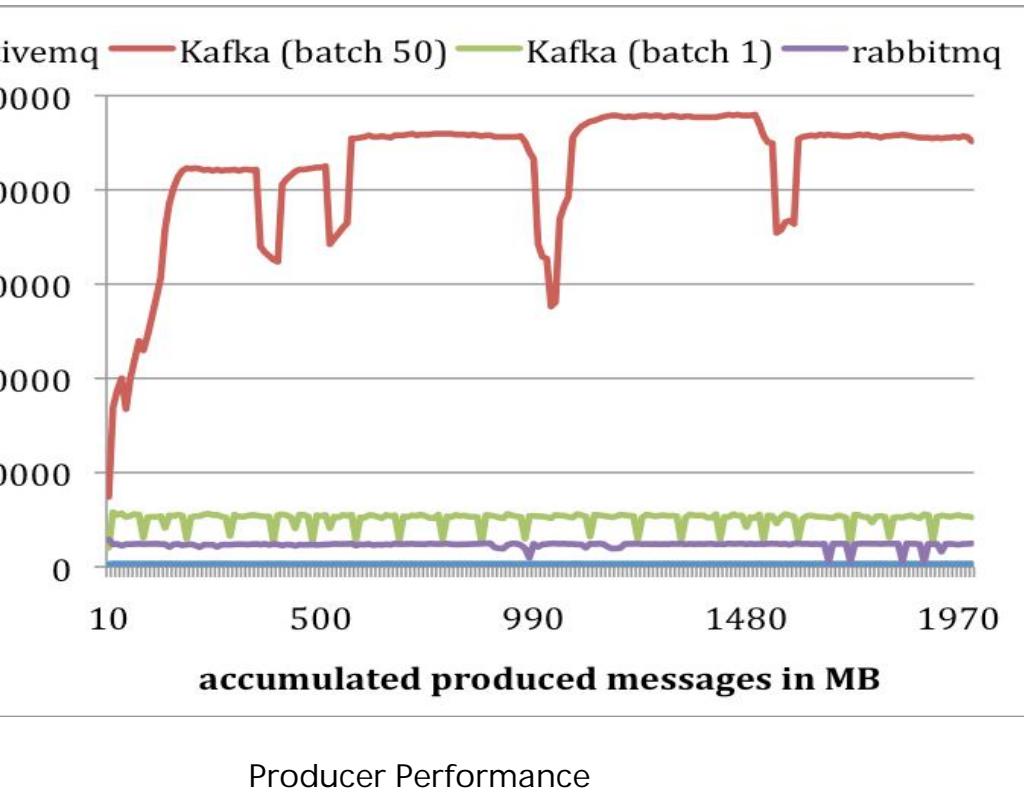
Send a record and invoke the given callback when the record has been acknowledged by the server

public Future<RecordMetadata> send(ProducerRecord<K, V> record, Callback callback);

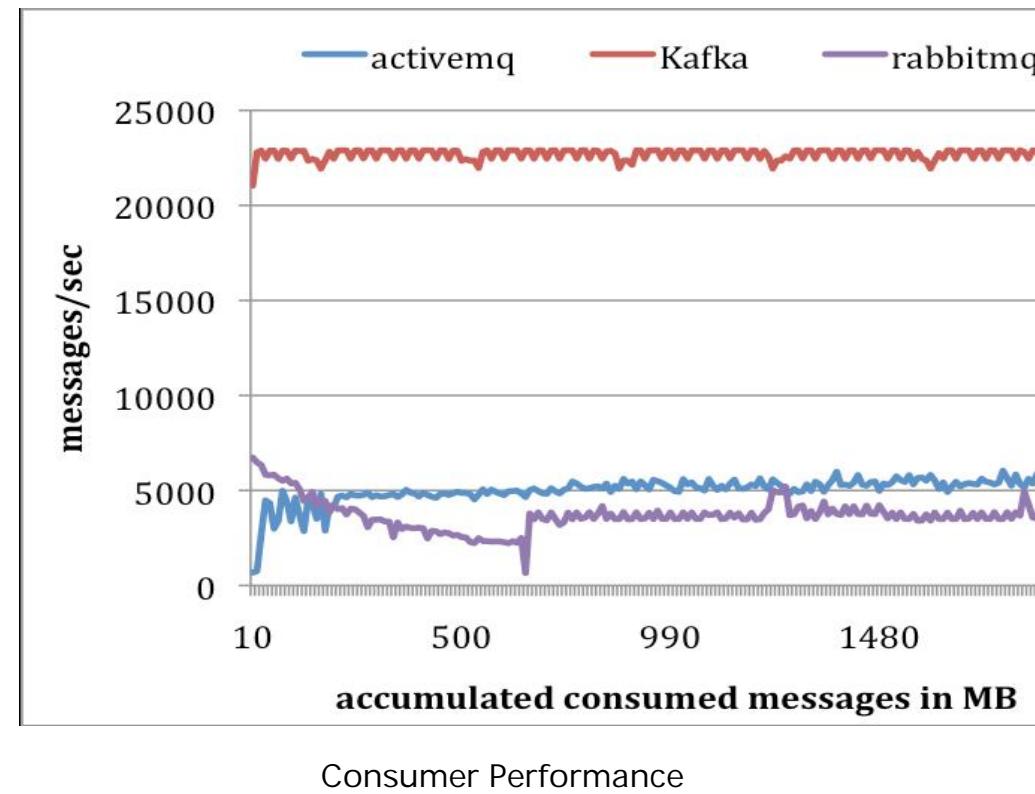
```
/ configure
Properties config = new Properties();
config.setProperty(ProducerConfig.BOOTSTRAP_SERVERS_CONFIG, "localhost:9092");
KafkaProducer producer = new KafkaProducer(config);

/ create and send a record
ProducerRecord<K, V> record = new ProducerRecord<K, V>("topic", "key".getBytes(), "value".getBytes());
Future<RecordMetadata> response = producer.send(record); // this is always non-blocking
System.out.println("The offset was: " + response.get().offset()); // get() blocks
```

Performance Numbers



Producer Performance



Consumer Performance

Credit : <http://research.microsoft.com/en-us/UM/people/srikanth/netdb11/netdb11papers/netdb11-final12.pdf>

Powered By

Apache Kafka is used at LinkedIn for activity stream data and operational metrics. This powers various products like LinkedIn Newsfeed, LinkedIn Today in addition to our offline analytics systems like Hadoop. [this](#)

[LinkedIn](#) - Apache kafka is used at LinkedIn as our main event bus that powers our news and activity feeds, automated review systems, and will soon power real time notifications and log distribution.

[Facebook](#) Kafka drives our new pub sub system which delivers real-time events for users in our latest game - Deckadence. It will soon be used in a host of new use cases including group chat and back end stats and log collection.

[Facebook](#) Apache Kafka aggregates high-flow message streams into a unified distributed pubsub service, brokering the data for other internal systems as part of Boundary's real-time network analytics infrastructure.

[Facebook](#) Kafka is used at DataSift as a collector of monitoring events and to track user's consumption of data streams in real time. <http://highscalability.com/blog/2011/11/29/datasift-architecture-realtime-datamining-at-120000-tweets>

[Facebook](#) We use Kafka to aggregate and process tracking data from all our facebook games (which are hosted at various providers) in a central location.

[Facebook](#) Kafka is used at AddThis to collect events generated by our data network and broker that data to our analytics clusters and real-time web analytics platform.

- At Urban Airship we use Kafka to buffer incoming data points from mobile devices for processing by our analytics infrastructure.
- We use Kafka to collect realtime event data from clients, as well as our own internal service metrics, that feed our interactive analytics dashboards.

We use Kafka internally as part of our reliable email queueing system.

We use a hierarchical distributed counting engine, uses Kafka as a primary speedy interface as well as routing events for cascading counting

We use Kafka to collect all metrics and events generated by the users of the website.

part of their Storm stream processing infrastructure, e.g. [this](#).

[this blog](#).

Kafka is part of the [InfoChimps real-time data platform](#).

[this blog](#).

[Foursquare](#) - We use Kafka as a distributed queue in front of our web traffic stream processing infrastructure (Storm).

[Foursquare](#) is used as the primary high speed message queue to power Storm and our real-time analytics/event ingestion pipelines.

[Foursquare](#) Kafka powers online to online messaging, and online to offline messaging at Foursquare. We integrate with monitoring, production systems, and our offline infrastructure, including hadoop.

[Foursquare](#) Kafka brokers data to most systems in our metrics and events ingestion pipeline. Different modules contribute and consume data from it, for streaming CEP (homegrown), persistence (at different "temperatures" in Redis, ElasticSearch, MySQL, etc), and search analysis (Hadoop).

[Twitter](#) We use Kafka 1. as an infrastructure that helps us bring continuously the tracking events from various datacenters into our central hadoop cluster for offline processing, 2. as a propagation path for data integration, 3. as a real-time processing layer for our recommendation engines

[Twitter](#) SPM (performance monitoring), Kafka is used for performance metrics collection and feeds SPM's in-memory data aggregation (OLAP cube creation) as well as our CEP/Alerts servers. In SA (search analytics) Kafka is used in several ways, as a source of data for indexing, as a source for search queries, and as a destination for search results before being aggregated and persisted in HBase.

[Twitter](#) We will soon be replacing part of our current production system to collect performance and usage data from the end-users browser for projects like Telemetry, Test Pilot, etc. Downstream consumers usually persist to either HDFS or MySQL.

[Waze](#) - At Waze Commerce (previously, NexTag), Kafka is used as a distributed queue in front of Storm based processing for search index generation. We plan to also use it for collecting user generated data on our web tier, landing pages, and mobile applications.

[Waze](#) We use Kafka as a bus to move all systems events through our various datacenters. This includes metrics, logs, custom events etc. On the consumer side, we output into Splunk, Graphite, Esper-like real-time alerting.

[Lucidworks](#) The Search Engine for Apps, Kafka is an integral part of our eventing, logging and messaging infrastructure.

[LinkSmart](#) Kafka is used at LinkSmart as an event stream feeding Hadoop and custom real time systems.

[Simple](#) We use Kafka at Simple for log aggregation and to power our analytics infrastructure.

[Lucidworks](#) - We use Kafka for syncing LucidWorks Search (Solr) with incoming data from Hadoop and also for sending LucidWorks Search logs back to Hadoop for analysis.

[Splunk](#) We use Kafka as a bus to move all systems events through our various datacenters. This includes metrics, logs, custom events etc. On the consumer side, we output into Splunk, Graphite, Esper-like real-time alerting.

Credit: <https://cwiki.apache.org/confluence/display/KAFKA/Powered+By>

[Hide Comments](#) | [Collapse All](#) | [Add Comment](#)